

Seungmin Oh 오승민

Ph.D. candidate in Artificial Intelligence

Multi-Modal Artificial Intelligence Lab, Ajou University, South Korea

Advised by Prof. Jongbin Ryu

seungmin.oh16@gmail.com

soonge.github.io

Scholar

github.com/SoongE

linkedin.com/in/soonge

Suwon, South Korea

Ph.D. candidate building **efficient AI** for deployment under tight compute limits, from edge devices to robot control loops, across vision, LLM, VLM, and VLA models.

Education

Integrated M.S./Ph.D. in Artificial Intelligence

Ajou University, Suwon, South Korea, advised by Prof. Jongbin Ryu

Feb 2022 - present

Ph.D. expected Feb 2028

B.S. in Software and Computer Engineering

Ajou University, Suwon, South Korea

Feb 2016 - Feb 2022

Research experience

Multi-Modal Artificial Intelligence Lab, Ajou University

Graduate Researcher, advised by Prof. Jongbin Ryu

Feb 2022 - present

Vision-language-action models for robotic manipulation

2025 - present

- Training methodology: world-state change prediction instead of raw actions, preserving VLM grounding (in prep.).
- Inference acceleration: fast action controller with lazy gating; full VLA invoked only on anticipated failure [S.3].
- Failure detection: early visual observer for semi-closed-loop action-chunk execution [S.1].

Efficient LLM compression

2024 - present

- Structured pruning: overcomplete training for scaled recovery capacity, up to +8.4 points at zero deployment cost [C.3].
- Structured pruning: curriculum-based layer-wise compression schedule, 1.7× faster convergence at half the GPU memory [C.4].

Efficient vision and test-time adaptation

2022 - present

- Test-time adaptation: candidate verification by divergence shift, up to 2× faster with over 3× less memory [C.5].
- On-device classification: 1st place, 2025 IEEE Low-Power Computer Vision Challenge Track 1 [A.1].
- Re-parameterization: branch-level neural substitution [C.1].
- Representation learning: transformation-aware prompt conditioning via re-calibrated contrastive loss [C.2].
- Medical: CT-based triage for small-bowel obstruction [J.1]; CNN vs. ViT architectural behavior [J.2].

Funded applied research projects (ETRI, NIA, 42dot)

2021 - 2024

- Real-time detection transformer for autonomous driving (42dot); source-free domain adaptation for on-device face recognition (ETRI); lymphoma recognition and survival analysis from 3D CT (NIA); masked-face recognition with domain adaptation (ETRI).

GGQ Company Inc., Seoul, South Korea

Research Intern

2020

- Game-scene understanding: player-name, skill, and HUD OCR in League of Legends; face recognition for attendance.

Publications

C conference, J journal, P patent, A preprint, S under review

Peer-reviewed publications and patents

- [C.5] **Seungmin Oh**, Seunghun Kang, Jongbin Ryu. Efficient Test-Time Adaptation through Candidate Verification and Divergence Shifts. *NeurIPS*, 2026.
- [C.4] Donggeon Lee, Dooyeon Na, **Seungmin Oh**, Jongbin Ryu. Layer-wise Curriculum Learning for Efficient LLM Compression. *EMNLP*, 2026.
- [C.3] **Seungmin Oh**, Donggeon Lee, Jongbin Ryu. Train Overcomplete, Deploy Compact: Scaling Recovery Capacity for Structured LLM Pruning. *EMNLP*, 2026.
- [C.2] **Seungmin Oh**, Seunghun Kang, Jongbin Ryu. Re-calibrated Contrastive Loss for Transformation-Aware Prompt Conditioning in Vision-Language Models. *BMVC*, 2026.
- [C.1] **Seungmin Oh**, Jongbin Ryu. Neural Substitution for Branch-Level Network Re-parameterization. *ACCV*, 2024.
- [J.2] **Seungmin Oh**, Namkug Kim, Jongbin Ryu. Analyzing to Discover Origins of CNNs and ViT Architectures in Medical Images. *Scientific Reports*, 14:8755, 2024.
- [J.1] **Seungmin Oh**, Jongbin Ryu, Hojung Shin, Jeongho Song, Sangyong Son, Hoon Hur, Sanguk Han. Deep Learning Using Computed Tomography to Identify High-Risk Patients for Acute Small Bowel Obstruction: Development and Validation of a Prediction Model. *International Journal of Surgery*, 109(12), 2023 (IF 15.3, JCR top 1%).

[P.2]

Jongbin Ryu, Dooyeon Na, **Seungmin Oh**, Donggeon Lee. System and Method for Compressing a Deep Learning Model Based on Layer-Wise Curriculum Learning. **Korean Patent Application No. 10-2026-0123632**, filed Jul 2026.

[P.1]

Hojung Shin, Jongbin Ryu, **Seungmin Oh**. Electronic Device and Method for Classifying Small Bowel Obstruction Using Abdominal Computed Tomography Images. **PCT International Patent Application No. PCT/KR2024/014604**, filed Sep 2024.

Preprints and submissions

[S.3]

Seungmin Oh, Jongbin Ryu. LazyVLA: Efficient Vision-Language-Action Inference through Controller-Guided Queries. **Under review**.

[S.2]

Yuri Kim, **Seungmin Oh**, Jongbin Ryu. Adaptive Regularization with Hardness-Aware Approach. **Under review**.

[S.1]

Donghyun Kim, **Seungmin Oh**, Dooyeon Na, Jongbin Ryu. Who Watches the VLA? Early Visual Observer for Semi-Closed-Loop Action-Chunked Execution. **Under review**.

[A.1]

Zihao Ye et al., including **Seungmin Oh**, Jongbin Ryu. Evaluation of Winning Solutions of 2025 Low Power Computer Vision Challenge [arXiv preprint](#), 2025.

Honors and awards

IEEE Low-Power Computer Vision Challenge, Winner CVPR 2025 Workshops, organized by IEEE, CVF, Qualcomm, and Purdue University; ranked 1st among 45 teams Layer Imitation: Efficient Vision-Language Model on the Mobile NPU link	Apr 2025
Google Solution Challenge, Global Top 100 Google 'Smiler,' an expression-training app for children with ASD link	May 2023
ICT Mentoring Conference, Excellence Award ACK, KOSIA Recommender System for Web Search Based on NLP to Improve User Search Environment link	Dec 2021
Student Paper Competition, First Prize Korea Digital Contents Society Lightweight Video Search Using Video-to-Text Generation and Retrieval	Jul 2021

Research funding

Doctoral Research Grant National Research Foundation of Korea (NRF), \$36K Progressive Linear Activation for Depth-wise Reparameterization of Large-scale Models	Sep 2025 - Aug 2027
Lesion-Class Re-labeling in Medical Image Training Data National Information Society Agency (NIA), \$20K	Oct 2022 - Dec 2023

Technical skills

Infrastructure	Built and operate a multi-node GPU cluster (about 20 nodes, 8× RTX 3090 each): containerized environment provisioning, per-node CUDA isolation, unattended long-running training pipelines.
Programming	Python; C/C++ with SIMD- and cache-aware performance optimization; CUDA.
ML tooling	PyTorch, Hugging Face, Weights & Biases, DVC, Docker, Kubernetes, Git and GitHub Actions.
On-device	Qualcomm AI Hub (Snapdragon NPU); structured pruning and re-parameterization for compact deployment.

Teaching and leadership

Lecturer , tutorials for undergraduate AI research assistants, Ajou University (3 terms)	2023 - 2024
Teaching Assistant , Capstone Design Project, Ajou University	2023
Founding Member , Google Developer Student Club, Ajou University chapter	Sep 2022 - Dec 2023